

CS342: Operating Systems Lab

Department of Computer Science and Engineering,

Indian Institute of Technology, Guwahati

North Guwahati, Assam 781 039

Exercise UP-01

OS Lessons: Process, Kernel data-structures, User and Kernel Stacks

Rating: Hard

Last update: 06 October 2017

This exercise is the first of the four exercises prepared to complete PintOS project User Programs. PintDoc lists this project as Project 2 and provides instructions and expectations in Chapter 3: User Programs.

Some general organizational arrangements and suggestions for the Exercise:

1. Use `cscope`, `ctags`, `gdb` and careful reading of the code to understand how the command made of a command-file name and its arguments is passed (or could be passed) from a calling function in the PintOS kernel code to the called function. Your search for the path that is necessary for setting up stack begins at function `run_action()` and ends at function `setup_stack()`. You may like to read the guidance for *Exercise UP-04* to get a better view of some specialized programming patterns that operating systems use to create the user programs as separate entities from the kernel codes.
2. The primary directory for project User Programs is `pintos/src/userprog`. To move to this directory, we need to run `make clean` in directory `pintos/src/threads`. Also, make the following changes to the primary script file.

Change Perl script `pintos` (which was previously placed in directory `~/bin`) as follows:

Line no. 24: Replace `"$HOME/pintos/src/threads/build/os.dsk"` by
`"$HOME/pintos/src/userprog/build/os.dsk"`
(This file was mentioned in PintOS installation instructions previously)

Further, please change the last line in file `pintos/src/userprog/Make.vars` to:
`SIMULATOR = --bochs`

[It is not needed in 2017 as `quem` has been installed]

1. This is also a good place to caution students about the limited amount of space available for the Kernel stack(s). To avoid stack overflow problems, do not write recursive functions in your kernel code.
2. In all modern operating systems, the kernel and the user virtual address spaces are separate, and all memory accesses across this boundary are carefully monitored through the features built into the processor hardware and kernel code.

36 Statement of the Exercise UP-01

37 A curt description of the tasks of the exercise is to write code to meet the specifications set in PintDoc
38 Section 3.3.3 *Argument Passing* on page 29. We will explain the specification in considerable detail
39 here as this exercise is challenging.

40 First, please read document [Executing main \(\) function on Linux](#) on the Internet.

41 The first dot-point listed in Section 3.2 (page 28) suggests that we must write code in function
42 `setup_stack()`. Section 3.1.4.1 on page 26 in PintDoc describes how the initial stack for
43 function `main()` of a PintOS user program is organized. You must also carefully read the example in
44 Section 3.5.1 *Program Startup Details* on pages 36-37.

45 We have already given a function to print the contents of this initial stack in a previous exercise. The
46 function is reproduced here (The code assumes that a c-pointer and `int` are same size objects.):

```
47 void test_stack(int *t)
48 {
49     int i;
50     int argc = t[1];
51     char ** argv;
52
53     argv = (char **) t[2];
54     printf("ARGC:%d  ARGV:%x\n", argc, (unsigned int)argv);
55     for (i = 0; i < argc; i++)
56         printf("Argv[%d] = %x pointing at %s\n",
57             i, (unsigned int)argv[i], argv[i]);
58 }
```

59 One important aim of this exercise can be stated as a question: Where can we call this function to
60 print the contents of the activation record that function `main()` will receive in the stack?

61 A search for the answer to the question will provide us a good overview of the kernel code that the
62 exercise must expand on the way to completing project User Programs.

63 User Interface to Create User Programs

64 There are two ways to start a user program in PintOS. One is to write a command line argument and
65 pass it to PintOS during the kernel load time. The kernel command-line directive `run` (see Section
66 3.1.2 *Using the File System*) takes a user command as its argument. The other way to start a user
67 program is through a system call from an already running user program; we will return to this method
68 in a later exercise (UP-04).

69 A user command may have arguments in addition to the name of a file containing code to run as a
70 user program. This makes some commands multi-word commands. A quote pair (" ") is used to
71 group such commands as a single argument for PintOS kernel load-time command directive `run`.
72 PintOS kernel load-time directives are similar to the built-in and separately-coded commands in
73 Unix/Linux shells. See near the bottom of page 24 [PintDoc] to learn how a command is specified as a
74 single argument to directive `run`. For this exercise, load-time commands are the only way we use to
75 run user programs.

76 What resources does a user program need to run?

77 A program running on a computer is called a *process*. A process is made of many tangible and virtual
78 components:

- 79 • A process needs a thread to execute program instructions and receive run-time on the
80 processor.
- 81 • Each process needs memory pages to store program instructions.
- 82 • Each process also needs initialized and uninitialized data segments in memory (also called
83 static data area). And,
- 84 • Each process needs a user stack to perform function calls (see Section 3.1.4.1 on page 26) for
85 passing arguments and return values.

86 A thread is ready to run as a user program/process when the above-listed resources are available to the
87 process. A thread, however, can run without some of these resources (components) as a kernel thread.

88 All ready-to-run threads are placed in `ready_list` and are scheduled (given time to run on the
89 computer processor) by PintOS scheduler. From the time the `run` directive receives a command to
90 run to the time the kernel has made available all resources needed by the process, the process is in a
91 state of *being prepared*.

92 Three Contexts and Three Phases

93 Students reading the kernel code for this exercise will find that the code is organized quite differently
94 from their previous understanding of program organizations. The reading of the code for this exercise
95 is easier if we view the code as grouped into three separate execution contexts.

96 The first context is the context of the thread that initiates the creation of a new user program as an
97 entity (process) to be run on PintOS operating system. We call this context as the context of the parent
98 thread. The kernel code that this parent thread runs to support the creation of the user program/process
99 makes an identifiable unit. A clear identification of this set of functions will help in reading,
100 understanding and modifying the kernel code as this code addresses some specific requirements. A
101 focus on these requirements would avoid distraction from the issues that concern the other parts of the
102 kernel code that run in other contexts.

103 The other contexts are the contexts of the thread created to execute user program code. This thread is
104 called a child thread. However, the contexts of a child thread are divided into two distinct phases or
105 contexts. We discuss the final phase first.

106 It is obvious that the third or the final context is the context in which the child thread is able to run the
107 user program. In this context, the child thread is primarily controlled by the instructions in the user
108 program and the thread's access to the kernel code is significantly limited and carefully regulated
109 through well-defined system calls. This, however, is the target context of the exercise – the very
110 purpose of this exercise.

111 The context of the execution between the first and the final context is the initial execution context of
112 the child thread. In this context, PintOS scheduler schedules the child thread and the thread runs the
113 kernel code. Once again, it will aid understanding if all functions accessed by the child thread in this
114 context are identified. This is so because the functions callable in this phase define a partition. This
115 containment will help you decide the changes you wish to implement for completing exercise UP-01.

116 There is another code partition in PintOS kernel code related to system calls. This is tiny in the
 117 provided PintOS code. Developing system calls will be the aim of the later exercises in this project.

118 Summary: It is important to understand the existence of three separate contexts of code execution as a
 119 kernel builds a user program. Three execution contexts of interest are (i) Context of the parent thread
 120 running kernel code, (ii) Context of the child thread running kernel code, and (iii) Context of the child
 121 thread running a user program. The kernel codes run in parent and child contexts are well separated in
 122 PintOS. The students should clearly identify the functions used in each of these two contexts.

123 There is strong *dependency* among the three contexts. The three contexts run in three phases one after
 124 another. This may cause some confusion to the naïve developers that the students doing this exercise
 125 obviously are. If a student focuses exclusively on the phases, they may not understand the association
 126 with the active thread running the code. Therefore, the synchronization issues may become difficult to
 127 comprehend and debug. Students are advised to keep the context of execution clearly in their view
 128 while developing the enhancements needed in this exercise.

129 A big advantage of building an understanding of the kernel code around contexts rather than phases is
 130 partitioning of the code. Phases do not partition the kernel code in obvious partitions.

131 **Where Does the Chicken Cross the Road?**

132 The three-context view of the kernel code provides a clear separation for writing code to complete this
 133 exercise. However, we have not yet guided the students to determine the groupings of the functions in
 134 PintOS code with each context.

135 User program codes are outside the kernel code. Only a little code in the provided PintOS kernel is
 136 part of the third (or User Program) context. Function `main()` of the user program is called when all
 137 needed resources for the user program process have been assembled. The creation of the child
 138 process is complete, and the child thread is ready for invoking function `main()` of the user program
 139 near label `done` in function `load()` in file `pintos/src/userprog/process.c`.

140 This obviously is the place to confirm that we have set the program stack correctly. Carefully
 141 determine where exactly you wish to call function `test_stack()` near label `done`.

142 The location in PintOS code that separates the parent thread's context from the child thread's context
 143 is obviously at the location where the child thread is added to `ready_list` and unblocked to run
 144 independently. Again, it is not too difficult to locate this in function `thread_create()` in file
 145 `threads/thread.c`. To provide further hints one notes that function `thread_create()` is
 146 called from function `process_execute()` in file `userprog/process.c`. An argument in the
 147 call lists `start_process()` as the function defining the starting point of the newly created child
 148 thread. This function is located in file `userprog/process.c`.

149 You notice that because the parent and the child threads are different threads, a tradition call to
 150 `start_process()` is not used in function `thread_create()`. If the threads running two
 151 functions were the same, then we would have expected a traditional call to function
 152 `start_process()` in function `process_execute()`.

153 It is unnecessary to say that all activities in function `start_process()` and functions called from
 154 it occur in the context of the child thread. Thus, the child thread is responsible for the loading of the
 155 user program and for setting up the initial stack for function `main()`. The thread is also responsible
 156 for invoking function `main()` and thus starting the new process. The thread must receive its

157 command – not just the (executable binary) file name – as a parameter to function
158 `start_process()`.

159 Before the Chicken Came on the Road

160 The responsibilities of the parent thread include the creation of the child thread and passing a single
161 command to the newly created child thread to run the intended user program. A command string is
162 made of the name of a program file containing an executable code and the arguments for the program.

163 The directive `run` is first noticed as a user command to load and run in kernel function
164 `run_actions()` of file `pintos/src/threads/init.c`. We need to understand the path
165 from `run_action()` to `thread_create()` in file `threads/thread.c`.

166 Since all these actions occur in the context of a single parent thread, the traditional practices of
167 program reading will be sufficient to trace the activities.

168 An appendix to this document lists a sequence of activities that helps in exploring the kernel code.

169 How we tested our implementation?

170 Once you have completed the exercise, we still have a small hurdle. Testing of the user program is not
171 possible yet! The reason for this limitation is non-availability of the system calls to write messages on
172 the computer console! Only kernel code can print messages; user programs cannot write on the
173 console yet.

174 We partially overcome this limitation by calling function `test_stack()` in the kernel code (and
175 not in the user program code). Caution: You may notice some differences in your output.

176 In the script below, the text typed by the user has been shown in bold. Some output from the standard
177 Pintos utilities has been deleted as it provides no useful insight. The output that is of minor interest is
178 shown in smaller fonts to fit the page width neatly.

179 The following commands were used to compile programs in directory `~/pintos/src/examples`

```
180 [vmm@progsrv ~]$ cd pintos/src/examples/  
181 [vmm@progsrv examples]$ make  
182 [Output deleted]
```

183 The following commands were used to setup Pintos to load and run a user program.

```
184 [vmm@progsrv examples]$ cd ../userprog/  
185 [vmm@progsrv userprog]$ make  
186 cd build && make all  
187 make[1]: Entering directory  
188 `/home/CS342/2016/FAC/vmm/pintos/src/userprog/build'  
189 make[1]: Nothing to be done for `all'.  
190 make[1]: Leaving directory  
191 `/home/CS342/2016/FAC/vmm/pintos/src/userprog/build'  
192 [vmm@progsrv userprog]$ cd build/  
193 [vmm@progsrv build]$ pintos-mkdisk fs.dsk 2  
194 [vmm@progsrv build]$ pintos -q -f  
195 [Output deleted]
```

```

196 Finally, we test our implementation of function stack_setup() :

197 [vmm@progsrv build]$ pintos -p ../../examples/echo -a echo -- -q
198 [vmm@progsrv build]$ pintos -q run "echo My stack_setup() works"
199 Writing command line to /tmp/EHglakwWBy.dsk...
200 squish-pty bochs -q
201 =====
202                      Bochs x86 Emulator 2.5.1
203                      Built from SVN snapshot on January 6, 2012
204                      Compiled on Oct 10 2012 at 11:12:02
205 =====
206 000000000000i[      ] reading configuration from bochsrc.txt
207 000000000000i[      ] installing nogui module as the Bochs GUI
208 000000000000i[      ] using log file bochsout.txt
209 Kernel command line: -q run 'echo My stack_setup() works'
210 Pintos booting with 4,096 kB RAM...
211 370 pages available in kernel pool.
212 369 pages available in user pool.
213 Calibrating timer... 204,600 loops/s.
214 hd0:0: detected 1,008 sector (504 kB) disk, model "Generic 1234", serial
215 "BXHD00011"
216 hd0:1: detected 4,032 sector (1 MB) disk, model "Generic 1234", serial
217 "BXHD00012"
218 Boot complete.
219 Executing 'echo My stack_setup() works':
220 ARGV:4  ARGV:bfffffc4
221 Argv[0] = bffffff0 pointing at echo
222 Argv[1] = bfffffed pointing at My
223 Argv[2] = bfffffdf pointing at stack_setup()
224 Argv[3] = bfffffd9 pointing at works
225
226 The last few lines above are of primary interest to verify the completion of the exercise. The
227 remaining output below is from Pintos code that has not yet been included in your project. You may
228 notice some variations in your output (but the variation is not relevant to your exercise.)

229 echo My stack_setup() works
230 echo: exit(0)
231 Execution of 'echo My stack_setup() works' complete.
232 Timer: 183 ticks
233 Thread: 0 idle ticks, 133 kernel ticks, 53 user ticks
234 hd0:0: 0 reads, 0 writes
235 hd0:1: 28 reads, 0 writes
236 Console: 819 characters output
237 Keyboard: 0 keys pressed
238 Exception: 0 page faults
239 Powering off...
240 =====
241 Bochs is exiting with the following message:
242 [UNMP ] Shutdown port: shutdown requested
243 =====
244 [vmm@progsrv build]$

```

Appendix

This appendix describes the story that traces the activities preceding the start of the execution of function `main()` in a user program on *Pintos*.

1. Makefiles provided in the project compile and link the *Pintos* kernel code and places the ready-to-load image of the kernel in the boot disk: `OS.dsk`.
2. You run programs on this kernel by using command pattern: `pintos arguments`. File `pintos` is a *Perl script* that interprets the command. String arguments are copied into the simulated disk `OS.dsk`. These arguments will be read by the kernel when it starts running.
3. Script `pintos` start running AI-32 simulator *Bochs*. *Pintos* kernel is loaded into this simulated computer.
4. Like any Unix program, *Pintos* kernel also starts its execution from function `main()`. *Pintos* function `main()` is in file `threads/init.c`.
5. The call that is of interest to understand the creation of a user thread is a call to function `run_actions(argv)` in file `threads/init.c`. Parameter `argv` carries the command line arguments we wrote to the script `pintos` in step 2.
6. This function in turn calls function `run_task()` in file `threads/init.c`. Here `argv` is split into individual tasks. A task is either a run of a user program or an in-built action.
7. From here the call (run request) goes to function `process_execute(task)` in file `userprog/process.c`. User program task will run as a child thread; the child thread will separate from the parent thread. (More on this in step 9 below)
8. The final function called by the parent thread is function `process_wait()` in file `userprog/process.c`. This function just terminates in the code provided in the initial implementation of *Pintos*.
9. However, before calling function `process_wait()` the parent thread calls function `process_execute(task)` in file `userprog/process.c`. Which calls function `thread_create` (with 4 arguments) in file `threads/thread.c`. This call creates a child thread and obliges the child to load the task code.
10. A child thread is created and allowed to run by calling `thread_unblock()`. This unblock is the formal point at which child and parent threads become separate entities and receive full access to processor time through *Pintos* scheduler.
11. The child thread begins its life at the start of function `start_process()` in file `userprog/process.c`. This function will load the user program and set up the stack for the call to function `main()` in the user program.
12. Child thread morphs into a user thread (or process) as it invokes and starts executing function `main()` of the user program.
13. You must also carefully study three (mysterious) stack frames that function `thread_create()` builds in memory allocated for `struct thread`.

Contributing Authors:

Vishv Malhotra, Gautam Barua, Rashmi Dutta Baruah